

FASHIONPRICEX: A CONTEXT-AWARE INTELLIGENT DYNAMIC PRICING FRAMEWORK FOR FASHION E-COMMERCE

Shobha Rani SH^{*1}, Yashaswini N.M^{*2}, Sowndarya AM^{*3}, Dr. A. Sathish^{*4}

^{*1,2}MCA (AI/DS), School Of Applied Science, Sapthagiri NPS University, Bengaluru, India.

^{*3}Assistant Professor, School Of Applied Science, Sapthagiri NPS University, Bengaluru, India

^{*4}Associate Professor, School Of Applied Science, Sapthagiri NPS University, Bengaluru, India.

DOI: <https://www.doi.org/10.56726/IRJMETS101575>

ABSTRACT

Pricing decisions in fashion e-commerce are unusually difficult because trends, seasons, sizes, and colours move quickly, unsold inventory loses commercial value with every week it sits on a shelf, and customer willingness to pay varies sharply across region, device, loyalty tier, and shopping occasion. Most production pricing systems still lean on static markdown calendars or simple rule-based discounting, which react slowly to shifting demand and leave revenue and margin on the table. This paper proposes FashionPriceX, a context-aware intelligent dynamic pricing framework that fuses customer context, product and inventory signals, competitor and market conditions, and temporal and seasonal patterns into a unified feature store feeding a demand-forecasting module and a reinforcement-learning-based pricing engine. A dedicated fairness and guardrail layer constrains the resulting price actions to stay within brand, legal, and ethical bounds before they reach a real-time pricing application programming interface. On a simulated fashion-retail environment built from synthetic transaction and browsing logs, FashionPriceX is compared against static pricing, rule-based discounting, and a regression-based machine-learning baseline. The proposed framework produces the largest simulated gains of the four strategies, improving revenue by 17.4%, gross margin by 11.2%, and reducing stockouts by 15.8% relative to static pricing. These results suggest that combining contextual personalisation with reinforcement-learning-driven price optimisation, under explicit fairness constraints, offers a practical path toward pricing in fashion e-commerce that is both more responsive and more accountable.

Keywords: Dynamic Pricing, Context-Aware Computing, Fashion E-Commerce, Reinforcement Learning, Demand Forecasting, Personalisation, Price Fairness.

I. INTRODUCTION

Fashion retail is one of the most volatile pricing environments in e-commerce. A single collection can span dozens of colourways, several size runs, and a selling window measured in weeks rather than months, and once a style falls out of trend its residual value drops sharply. Retailers therefore live with a persistent tension: price too high and inventory is marked down aggressively at season end, eroding margin; price too low and revenue is left on the table on styles that would have sold regardless. Because the underlying signals — trend momentum, size-curve depletion, weather, regional taste, and competitor promotions — change on a daily or even hourly basis, pricing decisions made once a week or once a season are structurally too slow for the market they serve.

The dominant pricing practice in fashion retail is still a markdown calendar: a style launches at full price and is discounted on a fixed schedule regardless of how demand is actually behaving. A more advanced but still limited alternative is rule-based discounting, where a small set of hand-written rules (for example, “cut price by 10% if sell-through is below target after two weeks”) approximates responsiveness without truly learning from data. Both approaches are transparent and easy to audit, but neither adapts to context: the same discount is applied to a loyal repeat customer and a first-time price-sensitive visitor, and to a region experiencing a heatwave and one entering winter, alike.

Machine learning has increasingly been proposed as a way to close this gap. Supervised models can forecast demand at a finer grain than manual rules allow, and reinforcement learning (RL) formulations treat pricing as a sequential decision problem in which an agent adjusts price and observes the resulting sales, gradually

learning a policy that balances immediate revenue against longer-run effects such as inventory depletion and customer trust. Field experiments on large marketplaces have reported measurable gains from this style of approach, and a growing body of research has explored Q-learning, deep Q-networks, and multi-agent formulations for retail and e-commerce pricing more broadly.

Fashion e-commerce adds a further requirement on top of demand forecasting and price optimisation: context-awareness. The same product can carry a different effective value for different customers and different moments — a customer who has browsed a category repeatedly without purchasing is not in the same position as one adding a final item to complete an outfit, and a jacket carries different urgency in a region entering a cold snap than in one experiencing unseasonal warmth. Incorporating this kind of contextual signal into pricing raises an additional concern that pure revenue-maximising systems tend to overlook: fairness. Unmitigated personalised pricing has been observed to produce meaningfully different prices for the same product shown to different customer segments, which creates reputational and, in some jurisdictions, regulatory risk if left unchecked.

This paper responds to these observations by proposing FashionPriceX, a context-aware intelligent dynamic pricing framework purpose-built for fashion e-commerce. The framework combines a context feature store, a demand-forecasting module, a reinforcement-learning-based pricing engine, and an explicit fairness and guardrail layer that constrains price actions before they are served to a storefront. On a simulated fashion-retail environment, FashionPriceX is benchmarked against static pricing, rule-based discounting, and a regression-based machine-learning baseline, and is shown to deliver the largest gains in revenue, margin, and stockout reduction among the strategies tested. The remainder of this paper reviews related work, describes the proposed methodology in detail, presents the simulated experimental results, and concludes with directions for future work.

II. LITERATURE SURVEY

Research on machine-learning-driven dynamic pricing for e-commerce has grown quickly over the last decade, and the resulting work can be grouped into four broadly overlapping strands: classical machine-learning approaches to price-setting, reinforcement-learning-based pricing agents, demand forecasting integrated with price optimisation, and context-aware personalisation together with the fairness questions it raises.

A first strand treats dynamic pricing as a supervised learning or classification problem. Nowak and Pawłowska-Nowak [1] proposed a dynamic pricing method for e-commerce that compares Naive Bayes, linear and nonlinear support vector machines, decision trees, and k-nearest neighbours on transaction data, finding that a linear support vector machine gave the most reliable classification of whether a price should be raised, lowered, or held. El Youbi, Messaoudi, and Loukili [3] took a regression-oriented view of the same problem, applying a gradient boosting machine to historical transaction data and reporting that boosted-tree ensembles captured pricing relationships that simpler models missed. Safonov [8] compared neural-network demand estimators against classical regression baselines and found that the neural approach generalised noticeably better on retail demand data. Together, this strand establishes that even relatively simple supervised models can meaningfully outperform static or purely rule-based pricing when enough transaction history is available, which motivates the demand-forecasting module built into FashionPriceX.

A second strand frames dynamic pricing as a sequential decision-making problem solved with reinforcement learning. Liu et al. [4] deployed a deep-reinforcement-learning pricing agent in a large-scale field experiment on Alibaba's Tmall platform and reported meaningful gains in both revenue and conversion relative to prices set manually by category managers. Apte, Kale, Datar, and Deshmukh [5] implemented a Q-learning agent in a simulated retail pricing environment and showed that the learned policy adapted prices to price-elasticity differences across products more effectively than a traditional constrained-optimisation baseline. Mussi et al. [7] extended this line of work to volume-discount pricing, introducing an online-learning algorithm evaluated on real product data from an e-commerce company and reporting a substantial increase in turnover relative to human-set discount schedules. These studies collectively support the choice of a reinforcement-learning pricing engine at the core of FashionPriceX, since RL formulations are naturally suited to a setting where the effect of a price change is only observed after the fact and must be learned through repeated interaction with the market.

A third strand integrates demand forecasting directly with price optimisation, which is especially relevant to fashion because so much of the assortment is newly introduced each season and has no direct sales history. Ferreira, Lee, and Simchi-Levi [2] developed a demand-forecasting and price-optimisation pipeline for Rue La La, an online fashion flash-sales retailer, using a non-parametric demand model together with a multi-product price-optimisation algorithm that accounts for reference-price effects; a field experiment reported a revenue increase of roughly ten percent without a corresponding drop in sales volume. Zhang and Zheng [12] studied optimal dynamic pricing under reference-price effects theoretically, showing that ignoring how customers anchor on previously seen prices can lead a naive pricing policy to under- or over-shoot the truly optimal price path. Yoshi et al. [6] combined structured transactional features with multilingual sentiment signals extracted from customer reviews in an XGBoost-based pricing model, finding that sentiment-aware features measurably improved prediction accuracy over models trained on transactional data alone. These results directly motivate the decision to give FashionPriceX its own demand-forecasting module rather than relying on the pricing engine to infer demand implicitly, and to model reference-price effects explicitly within the fairness and guardrail layer.

A fourth strand examines context-aware personalisation in recommendation and pricing, together with the fairness questions this raises. Zhao, Zhang, Friedman, and Tan [10] proposed predicting a customer's willingness-to-pay directly from purchase-decision boundaries and using that estimate to personalise discount depth at the point of recommendation. Kwon, Bak, and Bae [11] built a hybrid recommender that folds price similarity into collaborative filtering, reporting that explicitly modelling a customer's price sensitivity improved both revenue and click-through rate relative to a collaborative-filtering-only baseline in a large-scale online experiment. Wasilewski [9] proposed a context-aware quality framework for evaluating customer-segmentation algorithms used to personalise e-commerce interfaces, and separately, the wider personalisation literature has documented that unmitigated pricing and recommendation algorithms can produce systematically different outcomes for different demographic groups if fairness constraints are not built in from the start. This strand supports two design choices in FashionPriceX: first, that context should be modelled explicitly rather than left for the pricing engine to discover indirectly, and second, that a dedicated fairness and guardrail layer is necessary rather than optional.

Across all four strands, three gaps recur: most systems optimise price for a generic customer rather than a contextual one; demand forecasting and price optimisation are frequently treated as separate problems solved by separate teams; and explicit fairness constraints are rarely built into the pricing engine itself rather than checked after the fact. FashionPriceX is designed to address these three gaps directly by unifying context modelling, demand forecasting, reinforcement-learning-based price optimisation, and fairness guardrails within a single end-to-end framework, described in detail in the following section.

III. PROPOSED METHODOLOGY: THE FASHIONPRICEX FRAMEWORK

FashionPriceX is organised as a layered pipeline that moves from raw contextual data through to a governed, real-time price decision. The overall architecture is illustrated in Fig. 1, and its constituent stages are described in the subsections that follow.

A. Data Sources and Context Feature Engineering

FashionPriceX draws on four categories of context data. Customer context covers browsing and purchase history, loyalty tier, device type, and price sensitivity inferred from past behaviour. Product and inventory context covers style, colour, size-curve depletion, current markdown depth, and days since launch. Market and competitor context covers competitor pricing where observable, category-level demand momentum, and promotional activity. Temporal and seasonal context covers day of week, proximity to known sales events, and regional weather signals that are known to move demand for specific fashion categories, such as outerwear or swimwear.




Fig. 1. System architecture of the FashionPriceX dynamic pricing framework.

These four streams are cleaned, deduplicated, and encoded into a shared representation before being written to an online and offline feature store, from which the demand-forecasting and pricing modules read consistent, low-latency features at inference time. Fig. 2 summarises this pipeline in block-diagram form.




Fig. 2. Context feature engineering pipeline feeding the online and offline feature store.

B. Demand Forecasting Module

The demand-forecasting module estimates expected unit sales for a style at a candidate price, conditioned on the current context vector. It is implemented as an ensemble that combines a gradient-boosted tree model, which captures sharp, rule-like interactions such as size-curve exhaustion, with a recurrent (LSTM-based) component that captures smoother temporal trends such as trend momentum and seasonal build-up. The two sub-models are blended with weights that are periodically re-estimated on a rolling validation window, which lets the ensemble lean more heavily on whichever sub-model has been more accurate for a given category over the recent past. For newly launched styles with little or no sales history, the module falls back to a similarity-based estimate derived from comparable past styles matched on category, price point, and launch-season attributes, an approach broadly consistent with the reference-price and cold-start handling reported for fashion flash-sales pricing in prior work [2].

C. Reinforcement-Learning-Based Pricing Engine

The pricing engine treats each style-region combination as an episodic pricing task. At each decision step the agent observes the current context vector and the demand forecast, selects a price action from a discretised set of candidate multipliers on the base price, and receives a reward equal to the resulting margin-adjusted revenue, net of an inventory-holding penalty that discourages excessive overstock. The engine is implemented as a deep Q-network, chosen over tabular Q-learning because the context space is continuous and high-dimensional, and evaluated during development against a simpler Q-learning baseline consistent with prior retail pricing studies [5] to confirm that the added model capacity of the deep network translated into a measurable improvement on the simulated environment described in Section IV.

D. Fairness, Guardrail, and Business-Rule Layer

Before a price action reaches the storefront, it passes through a guardrail layer that enforces three kinds of constraints. Brand constraints cap the maximum discount depth and the maximum single-step price change for a given style, preventing abrupt or brand-damaging swings. Legal and policy constraints enforce minimum-advertised-price rules and jurisdiction-specific price-discrimination limits where applicable. Fairness constraints bound the maximum price difference the system is permitted to show for the same product across customer segments defined by protected or sensitive attributes, directly addressing the price-disparity risk documented in the personalisation literature discussed in Section II. Any price action that would violate a constraint is clipped to the nearest permissible value rather than rejected outright, so that the storefront always receives a valid price.

E. Real-Time Pricing API and Monitoring Dashboard

The guardrailed price action is exposed through a real-time pricing application programming interface that the storefront queries at page-render time, keeping end-to-end latency low enough for a normal browsing experience. A companion analytics and monitoring dashboard, illustrated later in Fig. 4, gives pricing analysts visibility into live price actions, aggregate revenue and margin trends, and any fairness-guardrail interventions, so that the system remains auditable rather than opaque.

IV. EXPERIMENTAL SETUP AND RESULTS**A. Simulated Experimental Environment**

Because a production fashion-retail deployment was outside the scope of this study, FashionPriceX was evaluated on a simulated environment built from synthetic transaction and browsing logs designed to reproduce realistic fashion-retail dynamics: category-specific price elasticities, seasonal demand curves, size-curve depletion, and randomly injected competitor promotions. The simulation spans twenty product categories and a twelve-month selling horizon, replayed under four pricing strategies so that every strategy is evaluated against an identical sequence of demand shocks.

B. Baselines

FashionPriceX is compared against three baselines. Static pricing holds the launch price fixed until a single season-end markdown. Rule-based discounting applies a small set of hand-written sell-through rules similar to

those used in current industry practice. Regression-based machine learning forecasts demand with a gradient-boosted model and chooses the revenue-maximising price given that forecast, without a reinforcement-learning feedback loop or a fairness and guardrail layer, and is included to isolate the specific contribution of the reinforcement-learning and guardrail components in the full FashionPriceX framework.

C. Results and Discussion

Fig. 3 summarises the simulated performance of the four strategies relative to static pricing. FashionPriceX achieved the largest gains on every metric tracked: a 17.4% increase in revenue, an 11.2% increase in gross margin, and a 15.8% reduction in stockouts, compared with 9.6%, 5.4%, and 8.5% respectively for the regression-based baseline, and smaller gains again for rule-based discounting.




Fig. 3. Simulated revenue, margin, and stockout improvement of each pricing strategy relative to static pricing on the fashion-retail test environment.

TABLE I. SIMULATED PERFORMANCE SUMMARY (RELATIVE TO STATIC PRICING)

Metric	FashionPriceX
Revenue Uplift	+17.4%
Gross Margin Uplift	+11.2%
Stockout Reduction	+15.8%
Fairness Guardrail Interventions	0 Unresolved

The gap between the regression-based baseline and full FashionPriceX suggests that the reinforcement-learning feedback loop, which learns from the downstream consequences of its own past price actions rather than only from a one-shot demand forecast, is responsible for a meaningful share of the improvement, consistent with the field-experiment gains reported for RL-based pricing agents elsewhere [4], [5], [7]. Throughout the simulation, the fairness and guardrail layer intervened to clip a small number of proposed price actions that would otherwise have produced discount depths or cross-segment price differences outside the configured bounds; no unresolved fairness violations remained in the final price actions served to the simulated storefront, which is summarised in Table I alongside the headline performance figures.

A live pricing-control dashboard, mocked up in Fig. 4, illustrates how these results would be surfaced to a pricing team in production: key performance indicators at a glance, a rolling revenue-index comparison against a baseline, per-category price actions, and the specific contextual signals currently driving the pricing engine's decisions.



Fig. 4. Mock-up of the FashionPriceX pricing-control and monitoring dashboard.

Taken together, these simulated results indicate that combining explicit context modelling, a dedicated demand-forecasting module, a reinforcement-learning pricing engine, and an enforced fairness layer yields larger and more balanced gains than any single technique applied on its own, while keeping the resulting price actions auditable and within defined ethical and business bounds.

V. CONCLUSION

This paper presented FashionPriceX, a context-aware intelligent dynamic pricing framework for fashion e-commerce that unifies customer, product, market, and temporal context within a shared feature store, feeds that context into a demand-forecasting module and a reinforcement-learning-based pricing engine, and constrains the resulting price actions through an explicit fairness and guardrail layer before they reach a real-time pricing application programming interface. On a simulated fashion-retail environment, the proposed framework outperformed static pricing, rule-based discounting, and a regression-based machine-learning baseline across revenue, margin, and stockout metrics, while keeping cross-segment price differences within defined fairness bounds.

Future work includes validating FashionPriceX on real transaction data from a production fashion retailer, extending the pricing engine to a multi-agent formulation that coordinates price actions across related styles and categories rather than treating each in isolation, and conducting a more formal audit of the fairness guardrail layer against a wider range of protected-attribute definitions and regional pricing regulations.

VI. REFERENCES

- [1] M. Nowak and M. Pawłowska-Nowak, "Dynamic Pricing Method in the E-Commerce Industry Using Machine Learning," *Appl. Sci.*, vol. 14, no. 24, art. 11668, 2024, doi: 10.3390/app142411668.
- [2] K. J. Ferreira, B. H. A. Lee, and D. Simchi-Levi, "Analytics for an Online Retailer: Demand Forecasting and Price Optimization," *Manuf. Serv. Oper. Manag.*, vol. 18, no. 1, pp. 69–88, 2016, doi: 10.1287/msom.2015.0561.
- [3] R. El Youbi, F. Messaoudi, and M. Loukili, "Machine Learning-Driven Dynamic Pricing Strategies in E-Commerce," in *Proc. 14th Int. Conf. Information and Communication Systems (ICICS)*, Irbid, Jordan, 2023, pp. 1–5, doi: 10.1109/ICICS60529.2023.10330541.
- [4] J. Liu, Y. Zhang, X. Wang, Y. Deng, and X. Wu, "Dynamic Pricing on E-Commerce Platform with Deep Reinforcement Learning: A Field Experiment," *arXiv:1912.02572*, 2020.

-
- [5] M. Apte, K. Kale, P. Datar, and P. R. Deshmukh, "Dynamic Retail Pricing via Q-Learning — A Reinforcement Learning Framework for Enhanced Revenue Management," arXiv:2411.18261, 2024.
- [6] A. M. Yoshi, A. Rohan, S. A. Mitu, M. M. K. Rabbi, S. Akther, and K. R. Ahmed, "Real-Time Dynamic Pricing Using Machine Learning: Integrating Customer Sentiment and Predictive Models for E-Commerce," Int. J. Adv. Comput. Sci. Appl., vol. 16, no. 9, pp. 22–30, 2025.
- [7] M. Mussi, G. Genalti, A. Nuara, F. Trovò, M. Restelli, and N. Gatti, "Dynamic Pricing with Volume Discounts in Online Settings," arXiv:2211.09612, 2022.
- [8] K. Safonov, "Neural Network Approach to Demand Estimation and Dynamic Pricing in Retail," arXiv:2412.00920, 2024.
- [9] A. Wasilewski, "Customer Segmentation in E-Commerce: A Context-Aware Quality Framework for Comparing Clustering Algorithms," J. Internet Serv. Appl., vol. 15, no. 1, pp. 160–178, 2024, doi: 10.5753/jisa.2024.3851.
- [10] Q. Zhao, Y. Zhang, D. Friedman, and F. Tan, "E-Commerce Recommendation with Personalized Promotion," in Proc. 9th ACM Conf. Recommender Systems (RecSys), 2015, pp. 219–226.
- [11] Y. Kwon, G. Bak, and Y. Bae, "A Hybrid Recommendation System Based on Similar-Price Content in a Large-Scale E-Commerce Environment," Appl. Sci., vol. 15, no. 19, art. 10758, 2025, doi: 10.3390/app151910758.
- [12] J. Zhang and Y. Zheng, "Optimal Dynamic Pricing for E-Commerce with Reference Effects," Manage. Sci., vol. 65, no. 11, pp. 5397–5415, 2019.
- [13] R. S. Sutton and A. G. Barto, Reinforcement Learning: An Introduction, 2nd ed. Cambridge, MA, USA: MIT Press, 2018.